

Data Analytics in Business

Classification and Regression Trees

(CART)

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Introduction to Tree Based Methods

Lessons in this Module

- A. Introduction to Tree Based Methods
- B. Regression Trees
- C. Regression Tree Examples
- D. Classification Trees
- E. Carseats Example for Classification Trees
- F. Bagging
- G. Random Forests
- H. Boosting
- I. Bagging, Random Forests, and Boosting Examples

* The materials in this module are from Chapter 8 of the ISLR textbook

Decision Trees: The Basics

- The main idea is that we segment the predictor space into simple regions.
- To make a prediction for an observation, find the region that the observation belongs to and then (typically) use the mean of the training data in that region as the predicted value.
- The set of splitting rules that are used to segment the predictor space can be summarized (and visualized) by a tree. Therefore, such approaches are also called **decision tree methods**.
- Decision tree methods are simple to implement and useful for interpretation.
- They can be used for **classification** or **regression** problems, hence they are now also known as **Classification and Regression Trees (CART)** methods.

GT_x

Decision Trees: The Basics

- We have a set of predictors X_1, X_2, \dots, X_p and we want to predict a response Y (quantitative or categorical value). We do this by growing a binary tree (i.e., each **internal node** has two branches).
- At each internal node of a tree, we apply a test to one of the predictors, say X_i
- Let's assume the test to be is $X_i < t$?
 - If this test is true (i.e., X_i 's value is $< t$) then we go to the left-hand branch.
 - If this test is false (i.e., X_i 's value is $\geq t$), then we go to the right-hand branch.
- Eventually, when we come to a **terminal** or **leaf node**, we make a **prediction**.
- This prediction will aggregate or average all the training data which reach that leaf node.

GT_x

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor
Scheller College of Business

Regression Trees

Regression Trees

- We use a simple example that uses the **Hitters** dataset in the ISLR package
- The objective is to predict a baseball player's Salary (measured in thousands of dollars) based on
 - Years (#of Years playing in the major leagues) and
 - Hits (# of Hits that he made the previous year)
- We clean up the data by dropping observations with missing Salary values.
- We also log-transform Salary so that its distribution is more normally-shaped.



GT_x

Regression Trees: *Hitters* Example

Figure 8.1 shows a regression tree for predicting the log salary of a baseball player.

- The split at the top node of the tree has two branches.
- The left-hand branch has observations for players with $\text{Years} < 4.5$, for which the mean log salary is 5.11. Therefore, we make a prediction in thousands of dollars, i.e., \$165,174
- How would you interpret the right branch of the tree?

Figure 8.1



GTx

Figure 8.2: The Three-Region Partition for the Hitters Data from Figure 8.1


The three regions can be written as

1. $R_1 = \{X \mid \text{Years} < 4.5\}$
2. $R_2 = \{X \mid \text{Years} \geq 4.5, \text{Hits} < 117.5\}$
3. $R_3 = \{X \mid \text{Years} \geq 4.5, \text{Hits} \geq 117.5\}$

Their respective predicted salaries are:

1. $\$1,000 * e^{5.107} = \$165,174$ for R_1
2. $\$1,000 * e^{5.999} = \$402,834$ for R_2
3. $\$1,000 * e^{6.740} = \$845,346$ for R_3

Figure 8.2




GTx

Terminology

- The regions R_1 , R_2 , and R_3 are known as **terminal nodes** or **leaves** of the tree.
- Decision trees are usually drawn upside down since the leaves are at the bottom of the tree.
- **Internal nodes** refer to the points in the tree where the predictor space is split.
- In Figure 8.1, the two internal nodes are Years < 4.5 and Hits < 117.5
- The segments of the trees that connect nodes are called **branches**.

Figure 8.1




GTx

How to Interpret Figure 8.1

- Years is the most important predictor in determining Salary. Players with less experience earn lower salaries than players with more experience.
- For a less experienced player, the number of hits made in the previous year seems to play little role in his salary.
- For players with equal to or more than 4.5 years of experience, the number of hits in the previous year does influence their salary, with players who made more hits tending to have higher salaries.
- The regression tree in Figure 8.1 is likely an over-simplification of the true relationship between Years, Hits, and Salary but is easier to interpret and has a helpful graphical representation.


Figure 8.1



GTx

How is a Regression Tree Built?

- How a regression tree is built:
 - We divide the predictor space (i.e., the set of possible values for the predictor variables) into J distinct and non-overlapping regions, R_1, \dots, R_J
 - For every observation that falls into a region R_j , we make the identical prediction, which is the mean of the response values for the training observations in that region.
- Using the example of Figure 8.1, for a given observation in region R_1 , we will predict a salary of \$165,174.
- The question remains – how are these regions – R_1, \dots, R_J – actually constructed?

GT_x

Constructing the Regions

- In theory, the regions could have any shape.
- However, for simplicity and ease of exposition of the resulting predictive model, we opt to divide the predictor space into high-dimensional rectangles or boxes.
- The goal is to find boxes R_1, \dots, R_J that minimize the SSE (Sum of Squared Errors)

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2,$$

where \hat{y}_{R_j} is the mean response for the training observations within the j th box

- Note that the ISLR book uses the term RSS (Residual Sum of Squares) instead of SSE. They are identical.

GT_x

Top-Down Greedy Approach

- Computationally, it is infeasible to consider every possible partition of the feature space into J boxes.
- Hence, we take a top-down, greedy approach called **recursive binary splitting**
 - Called top-down because it begins at the top of the tree with all observations in a single region and then successively splits the predictor space.
 - Every split has two new branches further down the tree.
 - It is greedy because, at each step of the process of tree-building, the best split is made at that particular step rather than looking ahead and picking a split that will lead to a better tree at a later step.

GTx

Recursive Binary Splitting

- Select the predictor X_j and the cutpoint s such that splitting the predictor space into the regions $\{X_j / X_j < s\}$ and $\{X_j / X_j \geq s\}$ leads to the greatest possible reduction in SSE.
- We have to consider all predictors X_1, X_2, \dots, X_p and all possible values of s for each of the predictors, and then choose the predictor and its cutpoint such that the resulting tree has the lowest SSE. This means that for any j and s , we define the pair of half-planes $R_1(j, s) = \{X_j / X_j < s\}$ and $R_2(j, s) = \{X_j / X_j \geq s\}$, and we seek the value of j and s that minimize the equation

$$\sum_{i; x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i; x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

where \hat{y}_{R_1} is the mean response for the training observations in $R_1(j, s)$ and

where \hat{y}_{R_2} is the mean response for the training observations in $R_2(j, s)$

- Finding the values of j and s that minimize the equation can be done quickly, especially when p is not too large.

GTx

Recursive Binary Splitting

- We repeat the process looking for the best predictor and cutpoint to split the data further so as to minimize the SSE within each of the split regions.
- Note that we are now splitting one of the two previously identified regions so we now have three regions.
- Again, we look to split one of these three regions further, so as to minimize SSE.
- The process continues until a stopping criterion is reached; for example, we could continue until no region has more than five observations.
- Once the regions R_1, \dots, R_J have been created, we predict the response for a given test observation using the means of the training observation in the region in which that test observation belongs.

GT_x

Tree Pruning

- Recursive binary splitting may produce good predictions on the training set but is likely to overfit the data, leading to poor performance on the test set. This is because the tree might be too complex.
- A smaller tree with fewer splits (i.e., fewer regions R_1, \dots, R_J) might lead to
 - lower variance
 - better interpretation at the cost of a little bias
- One alternative is to build a tree as long as the decrease in SSE due to each split exceeds a (high) threshold.
- This will result in smaller trees but could be short-sighted since an initial seemingly-worthless split might be followed by a very good split later on.

GT_x

Pruning to Build a Better Tree

- A better strategy is to grow a very large tree (call it T_0) and then prune in back to obtain a subtree.
- How does one find the best subtree?
 - Our aim is to find a subtree that leads to the lowest test error rate.
 - Given a subtree, we can estimate the test error rate using the validation set approach or using cross-validation (CV).
 - But estimating the CV for every possible subtree would take a long time since there are so many subtrees.
 - So we need a way to select a small set of subtrees to consider.

GTx

Cost Complexity Pruning

- **Cost complexity pruning** or **weakest link pruning** is a way to select a small set of subtrees to consider.
- Instead of examining every possible subtree, we consider a sequence of trees indexed by a nonnegative tuning parameter α
- $|T|$ indicates the number of terminal nodes of a tree T
- R_m is the rectangle or box corresponding to the m th terminal node
- \hat{y}_{R_m} is the mean response for the training observations within the m th box
- α controls a trade-off between the subtree's complexity and its fit to the training data and for each value of α there is a subtree $T \in T_0$, such that

$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T| \quad \text{Equation 8.4}$$

is as small as possible.

GTx

Cost Complexity Pruning

Algorithm 8.1 *Building a Regression Tree*

1. Use recursive binary splitting to grow a large tree on the training data, stopping only when each terminal node has fewer than some minimum number of observations.
 2. Apply cost complexity pruning to the large tree in order to obtain a sequence of best subtrees, as a function of α .
 3. Use K-fold cross-validation to choose α . That is, divide the training observations into K folds. For each $k = 1, \dots, K$:
 - (a) Repeat Steps 1 and 2 on all but the k th fold of the training data.
 - (b) Evaluate the mean squared prediction error on the data in the left-out k th fold, as a function of α .

Average the results for each value of α , and pick α to minimize the average error.
 4. Return the subtree from Step 2 that corresponds to the chosen value of α .
-

* Chapter 8 of the ISLR textbook

GT_x

Notes on the Algorithm

$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T| \quad \text{Equation 8.4}$$

- When $\alpha = 0$, then the subtree T will simply be equal to T_0 , because Equation 8.4 measures the training error.
- As α increases, there is a price to pay for having a tree with many terminal nodes and so Equation 8.4 will be minimized when we have a smaller tree.
- Note that Equation 8.4 is similar to the Lasso which was used to control the complexity of the linear model.
- As we increase α from 0 in Equation 8.4, branches get pruned from the tree in a nested and predictable manner, which means obtaining the whole sequence of subtrees as a function of α is easy.

GT_x

Quiz

A regression tree can approximate only linear relationships, as each split is linear.

- A. True
- B. False (overall boundary after second split is non-linear)
- C. Depends on how we prune
- D. Depends on which variable is chosen first

Correct Answer: B

Which of these is a FALSE statement about Decision Trees?

- A. At each internal node of a tree, we apply a test to one of the predictors, say X_i
- B. When we come to a leaf node, we make a prediction
- C. Decision trees are usually drawn with the leaves at the top of the tree
- D. Recursive binary splitting may produce good predictions on the training set but is likely to overfit the data, leading to poor performance on the test set
- E. Every split has two new branches further down the tree

Correct Answer: C

GTx

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Regression Tree Examples

Regression Tree for the *Hitters* Data Set

- There are several R packages for regression trees
 - `tree()` and `rpart()` are two packages
- Remove observations with missing data (`Hitters` or `Salary`)
- Log-transform `Salary` so that its distribution is closer to normal
- Then build a tree

```
library(ISLR)
library(tree)
attach(Hitters)
# remove NA values
Hitters <- na.omit(Hitters)
Salary <- na.omit(Salary)
# put salary on log scale and fit reg. tree
treefit <- tree(log(Salary) ~ Years + Hits, data=Hitters)
```

GT_x

Summary of Fitted Tree

```
summary(treefit)
##
## Regression tree:
## tree(formula = log(Salary) ~ Years + Hits,
## data = Hitters)
## Number of terminal nodes: 8
## Residual mean deviance: 0.2708 = 69.06 / 255
## Distribution of residuals:
##   Min.   1st Qu. Median Mean   3rd Qu.  Max.
## -2.2400 -0.2980 -0.0365  0.0000  0.3233  2.1520
```


Do a summary of `treefit` – the fitted regression tree.

There are 8 terminal nodes or leaves of the tree. Here “deviance” is just the sum of squared errors; this gives us a $MSE = 69.056 / (263 - 8) = 69.06 / 255 = 0.2708$

GT_x

Plot the Tree

```
# plot the regression tree.
plot(treefit)
text(treefit,cex=0.75)
# note leaf labels are log salary!
```



GTx

Another Regression Tree Example

```
### Regression Tree using the Boston data set in MASS library
```

```
set.seed (1)
train = sample(1:nrow(Boston), nrow(Boston)/2)
tree.boston=tree(medv~.,Boston, subset=train)
summary (tree.boston)
```

GTx

Cross Validation and Pruning

- The tree package contains functions **prune.tree()** and **cv.tree()** for pruning trees by cross-validation.
- **prune.tree()** uses as input a tree that you have fitted by **tree()** and determines a nested sequence of subtrees of the supplied tree by recursively “snipping” off the least important splits.
- Evaluation can be done on new data or on the training data.
- The default is method=“deviance”, which fits by minimizing the mean squared error (for continuous responses) or the negative log likelihood (for discrete responses).
- **cv.tree()** runs a K-fold cross-validation experiment to find the deviance or number of misclassifications as a function of the cost-complexity parameter.


GTx

Cross Validation and Pruning

use the **cv.tree()** function to see if pruning will improve performance

```
cv.boston = cv.tree(tree.boston)
plot(cv.boston$size, cv.boston$dev,
     type='b')
```

- The most complex tree is selected in this case by CV




GTx

Cross Validation and Pruning

- But, if we wish to prune the tree, we could do so as follows, using the `prune.tree()` function:

```
prune.boston
=prune.tree(tree.boston,best=5)
plot(prune.boston)
text(prune.boston,pretty =0)
```

GT_x

Cross Validation and Pruning

- # In keeping with the cross-validation results, we use the # unpruned tree to make predictions on the test set:

```
yhat=predict(tree.boston, newdata=Boston[-train,])
boston.test=Boston[-train ,"medv"]
plot(yhat,boston.test)
abline(0,1)
mean((yhat - boston.test)^2)
# [1] 15.26222
```

- In other words, the test set MSE associated with the regression tree is 25.05. The square root of the MSE is therefore around 5.005, indicating that this model leads to test predictions that are within around \$5,005 of the true median home value for the suburb.

GT_x

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor
Scheller College of Business

Classification Trees

Classification Trees

- Classification trees are very similar to regression trees.
- The difference is that the response (Y variable) is qualitative rather than quantitative.
- For a classification tree, the predicted response for an observation is to put it in the most commonly occurring class of training observations in the region to which that observation belongs.
- We want to know how to interpret the results of a classification tree

Growing a Classification Tree

- The process of growing a classification tree is very similar to that of growing a regression tree.
- Recursive binary split is used to grow the tree.
- However, we can't use the SSE as a criterion for making splits.
- **Classification error rate** is used as the criterion for making splits.
- The classification error rate is defined as the fraction of the training observations in a region that do not belong to that region's most common class:

$$E = 1 - \max_k \hat{p}_{mk}$$

where \hat{p}_{mk} represents the proportion of training observations in region m that belong to class k

- However, the classification error is not sufficiently sensitive to the process of growing a tree. Two other measures will be discussed next.

GTx

The Gini Index

- The **Gini index** is measure of the total variance across all the K classes (in the m th region) and is defined as:

$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

- The Gini index has a very small value when all of the \hat{p}_{mk} 's are close to 0 or 1.
- It is also referred to as a measure of **node purity**. A small value of G indicates that a node mostly contains observations from a single class.

GTx

Entropy

- An alternative measure to the Gini index is **entropy**, which is defined as:

$$D = - \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$$

- Since $0 \leq \hat{p}_{mk} \leq 1$, we get $0 \leq -\hat{p}_{mk} \log \hat{p}_{mk}$
- Entropy will have a very small value if all of the \hat{p}_{mk} 's are close to 0 or 1.
- This means that entropy will also be close to 0 if the m th node is pure.
- As it happens, the Gini index and entropy are very similar numerically.

GT_x

Using These Measures

- Either the Gini index or Entropy is used to evaluate the quality of a particular split when building a classification tree since they are very sensitive to purity compared to the classification error rate.
- They could also be used for pruning a tree, but the classification error rate is preferable if prediction accuracy of the final pruned tree is desired.

GT_x

Deviance

- When one uses the `summary()` function on a classification tree, **deviance** is reported which is defined as:


$$-2 \sum_m \sum_k n_{mk} \log \hat{p}_{mk}$$

where n_{mk} is the number of observations.

- A small deviance is a sign that the tree provided good fit to the (training) data.
- The residual mean deviance reported is equal to the deviance divided by $n - |T_0|$

GTx

Quiz



Is a tree based method suitable for the above two-dimensional classification problem where each region is shaded a different color (see figure above)?

- Yes
- No, a linear model is more suitable
- Depends on the tree that forms

Correct Answer: B

GTx

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor
Scheller College of Business

**Carseats Example for
Classification Trees**

Carseats Dataset

- In the Carseats dataset in the ISLR package, Sales is a continuous variable and so we create a new binary variable (**High**) to build a classification tree

```
library(ISLR)
library(tree)
attach(Carseats)
# creating a binary variable
High <- ifelse(Sales <= 8, "No", "Yes")
# merge High with the rest of the Carseats
Carseats <- data.frame(Carseats, High)
```

- Note that 236 out of 400 records have Sales ≤ 8 (so High is set to No for them) and the remaining 164 have Sales > 8 (so High is set to Yes for them)

GT_x

Fit a Classification Tree

- Use the `tree()` function to fit a classification tree to predict **High** using all variables but Sales

```
a <- tree(formula = High ~ . - Sales, data = Carseats)
summary(a)
```

Classification tree:

```
tree(formula = High ~ . - Sales, data = Carseats)
```

Variables actually used in tree construction:

```
[1] "ShelveLoc" "Price" "Income" "CompPrice" "Population"
"Advertising"
[7] "Age" "US"
```

Number of terminal nodes: 27 # = $|T_\theta|$

Residual mean deviance: 0.4575 = 170.7 / 373 # (= deviance / (n - $|T_\theta|$))

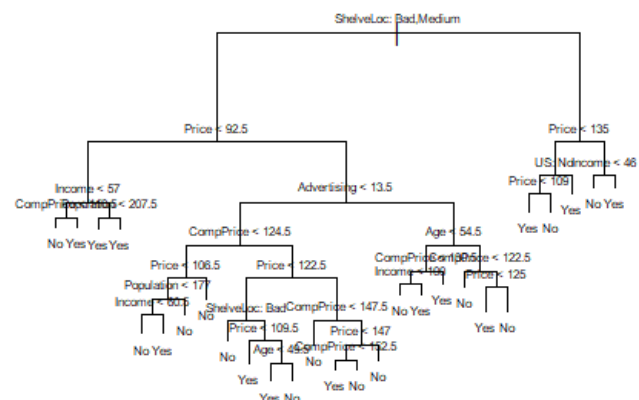
Misclassification error rate: 0.09 = 36 / 400 # (= training error rate)

GT_x

Plot the Tree

```
plot(a)
text(a, cex=0.60,
pretty=0)
```

- The most important indicator of Sales seems to be Shelf location since the 1st branch separates **good** locations (on the right) from **bad** or **medium** ones (on the left).



Type the Name of the Tree

- If we type the name of the tree object,
 - R prints the output corresponding to each branch of the tree,
 - It displays the split criterion,
 - # of observations in that branch,
 - the deviance,
 - overall prediction for that branch (Yes or No),
 - and the fraction of observations in that that take the value of Yes or No.
 - Branches that lead to terminal nodes are marked with an asterisk *

GT_x

The Classification Tree a First Few Lines of Output

```
> a
node), split, n, deviance, yval, (yprob)
  * denotes terminal node

1) root 400 541.500 No ( 0.59000 0.41000 )
 2) ShelfLoc: Bad,Medium 315 390.600 No ( 0.68889 0.31111 )
  4) Price < 92.5 46 56.530 Yes ( 0.30435 0.69565 )
    8) Income < 57 10 12.220 No ( 0.70000 0.30000 )
      16) CompPrice < 110.5 5 0.000 No ( 1.00000 0.00000 ) *
      17) CompPrice > 110.5 5 6.730 Yes ( 0.40000 0.60000 ) *
    9) Income > 57 36 35.470 Yes ( 0.19444 0.80556 )
      18) Population < 207.5 16 21.170 Yes ( 0.37500 0.62500 ) *
      19) Population > 207.5 20 7.941 Yes ( 0.05000 0.95000 ) *
```

at the root node 236/400 = 0.59 of the records have sales <= 8 (High = 0), hence the predicted value of Y = 0 (i.e., High = No)

nodes 16, 17, 18, 19 are terminal nodes

GT_x

Estimating the Test Error to Evaluate Tree Performance

- Split the observations into a training set and a test data.
- Build the tree using the training set.
- Evaluate the tree's performance on the test set using the predict() function.
- For a classification tree, the argument type="class" tells R to return the actual class prediction.

GT_x

Prediction Using Test Data

```
# prediction performance on validation set
set.seed(2)
train=sample(1:nrow(Carseats), 200)
Carseats.test=Carseats[-train,]
High.test=High[-train] # actual High values on the observations in the test set
tree.carseats=tree(High ~ . - Sales, data = Carseats, subset=train)
tree.pred=predict(tree.carseats, Carseats.test, type="class")
table(tree.pred,High.test)
      High.test
tree.pred No  Yes
      No   86   27
      Yes  30   57
(86+57)/200
[1] 0.715
```

The prediction accuracy is 71.5% for observations in the test data set

GT_x

Pruning the Tree to Try to Improve Prediction

- Let's prune the tree to see if we get better prediction results.
- The function `cv.tree()` performs cross-validation to determine the optimal level of tree complexity; cost complexity pruning is how the sequence of trees are considered.
- The argument `FUN=prune.misclass` indicates that we want the classification tree error rate to guide the CV and pruning process rather than the default for the `cv.tree()` function which is the deviance.
- The `cv.tree()` function reports the # of terminal nodes of each tree considered (size), as well as the corresponding error rate and the value of the cost complexity parameter used (k , which corresponds to α in equation 8.4).

GTx

Cross Validation Using `cv.tree`

```
# cross validation
set.seed(3)
cv.carseats=cv.tree(tree.carseats,FUN=prune.misclass)
names(cv.carseats)
cv.carseats
```

GTx

R Output of cv



```
> cv.carseats
$`size`
[1] 19 17 14 13 9 7 3 2 1
$dev
[1] 55 55 53 52 50 56 69 65 80
$k
[1] -Inf 0.0000000 0.6666667 1.0000000 1.7500000 2.0000000
4.2500000 5.0000000
[9] 23.0000000
$method
[1] "misclass"
attr(,"class")
[1] "prune" "tree.sequence"
```

dev corresponds to the cross-validation error rate in this example
 # size corresponds to # of terminal nodes
 # tree with 9 terminal nodes has lowest cv error rate of 50

GTx

Plot Error Rate as a Function of Size and k

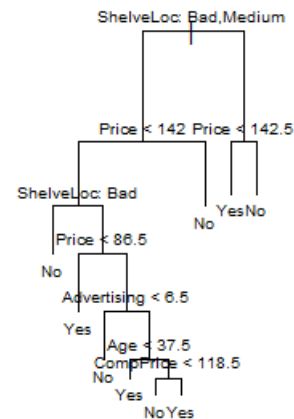
```
par(mfrow=c(1,2))
plot(cv.carseats$size,
cv.carseats$dev,type="b")
plot(cv.carseats$k,
cv.carseats$dev,type="b")
```



GTx

9-Node Tree Obtained from Pruning

```
# 9-node tree from pruning
prune.carseats=prune.misclass
(tree.carseats,best=9)
plot(prune.carseats)
text(prune.carseats,pretty=0,
cex=0.6)
```



GTx

Prediction Accuracy of the 9-Node Tree

```
tree.pred=predict(tree.carseats, Carseats.test,type="class")
table(tree.pred,High.test)
```

```
> table(tree.pred,High.test)
      High.test
tree.pred No  Yes
      No   87   26
      Yes  29   58
> (87+58)/200 # this is the accuracy.
[1] 0.725
```

Pruning has produced a more interpretable tree with improved prediction accuracy

GTx

Prediction Accuracy of the 15-Node Tree

- If we increase the value of best to 15, we get a larger pruned tree with lower prediction accuracy.
- Verify this on your own!
- # 15-node tree from pruning

```
prune.carseats=prune.misclass(tree.carseats,best=15)
plot(prune.carseats)
text(prune.carseats,pretty=0,cex=0.6)
# prediction performance of this 15 node tree
tree.pred=predict(tree.carseats, Carseats.test,type="class")
table(tree.pred,High.test)
```

GT_x

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Bagging

Bagging

- When we use decision trees for prediction, they could suffer from **high variance**. If we split the training data into two parts at random, and fit a decision tree to both halves, the results we get could be quite different.
- On the other hand, a procedure with low variance will give similar results if applied repeatedly to distinct data sets; linear regression tends to have low variance if n is moderately larger than p
- **Bootstrap aggregation** or **bagging** is general-purpose technique for reducing the variance of a statistical learning method; it is particularly useful in the context of decision trees.

GTx

Bagging

- If we are given a set of independent observations Z_1, \dots, Z_n each with variance σ^2 , the variance of the mean \bar{Z} of these observations is given by σ^2/n
- So, averaging a set of observations reduces variance.
- So, a natural way to reduce the variance and thus increase the prediction accuracy of a statistical learning method, is to
 - take many training sets from the population,
 - build a separate model using each training set,
 - and average the resulting predictions.

GTx

Bagging

- We calculate B predicted values $f^1(x), f^2(x), \dots, f^B(x)$ using B separate training sets.
- We average them to get a single low-variance statistical learning model

$$\hat{f}_{avg}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$$

- However, doing this process isn't practical because we mostly don't have multiple training sets.

GTx

Bagging

- But bootstrap gives us a workaround.
- We can take repeated samples from a single training data set.
- We then generate B different bootstrapped training sets.
- The learning method is trained on the b th bootstrapped training set to obtain the prediction $\hat{f}^{*b}(x)$

- We average all B predictions to get

$$\hat{f}_{bag} = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

- This technique is called **bagging**.

GTx

Bagging and Regression Trees

- For regression trees, the bagging procedure works by building B regression trees using B bootstrapped training sets.
- The predictions of these B regression trees are then averaged.
- These trees are grown deep and not pruned.
- Each such tree has high variance but has low bias.
- Averaging the predictions of these B trees reduces variance.
- Bagging has been shown to give impressive accuracy improvements by combining together hundreds or thousands of trees into one procedure.

GT_x

Bagging and Classification Trees

- The simplest way to apply bagging to classification trees is, for each test observation,
 - record the class predicted by each of the B trees for that observation, and
 - then take a majority vote
- In this context, a majority vote is the overall prediction of the most commonly occurring class among all B predictions for that test observation.

GT_x

Out-of-Bag Estimation

- To estimate the test error of a bagged model, there is a straightforward way without using cross-validation or validation set approach.
- In bagging, as we have seen, trees are repeatedly fit to bootstrapped subsets of the observations.
- One can show that, on average, each bagged tree makes use of around $2/3$ of the observations.
- The remaining $1/3$ of the observations that are not used to fit a given bagged tree are called **out-of-bag** (OOB) observations.
- We can predict the response for the i th observation using each of the trees in which that observation was OOB.
- This will yield around $B/3$ predictions for that i th observation.

GTx

Out-of-Bag Estimation

- To obtain a single prediction for the i th observation, we can average these predicted responses (for regression trees) or take a majority vote (for classification trees).
- This gives a single OOB prediction for the i th observation
- Do this for all n observations and then compute the overall OOB MSE (for regression tree) or classification error (for classification tree).
- This resulting OOB error is a valid estimate of the test error for the bagged model because the response for each observation is predicted using only those trees that were fit not using that observation.
- It can be shown that with B sufficiently large, OOB error is virtually similar to LOOCV.
- This OOB approach is very convenient for bagging on large data sets where cross-validation would be computationally burdensome.

GTx

Variable Importance Measures

- We saw that bagging typically yields improved accuracy over prediction that only uses a single tree but it is difficult to interpret the model that results from bagging.
- It not possible to represent the model using a single tree because it is no longer clear which variables are most important to the procedure.
- Thus the improved prediction accuracy comes at the cost of interpretability but one can obtain an overall summary of the importance of each predictor using SSE (for bagging regression trees) or the Gini index (for bagging classification trees)
- When bagging regression trees, we can record the total amount the SSE is reduced due to splits over a given predictor, averaged over all B trees. A large average value indicates an important predictor
- Similarly, when bagging classification trees, we can add the total amount that the Gini index is decreased by splits over a given predictor, averaged over all B trees.

GT_x

Quiz

When bagging regression trees, we can record the total amount the SSE is reduced due to splits over a given predictor, averaged over all B trees. A large average value indicates an important predictor

- A. TRUE
- B. FALSE
- C. MAYBE

Correct Answer: A

GT_x

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor
Scheller College of Business

Random Forests

Random Forests

- Random Forests are an improvement over bagged trees by way of a small tweak that **decorrelates** the trees.
- Similar to bagging, we build a number of trees on bootstrapped training samples.
- But in Random Forests, when building the decision trees, each time a split is considered, **a random sample of m predictors** is chosen from the full set of p predictors
- The split is allowed to use only of these m predictors.
- A fresh sample of m predictors is used at each split, and typically $m \approx \sqrt{p}$
- Note that in building a random forest, the algorithm is not even allowed to consider a majority of the p predictors.
- This has a good rationale.

GTx

Random Forests

- Assume that there is one very strong predictor in the data set along with a number of other moderately strong predictors.
- Then in the collection of bagged trees, most or all of the trees will use this strong predictor in the top split.
- Therefore, all of the bagged trees will look similar to each other and thus their predictions will be highly correlated.
- Note that averaging many highly correlated quantities does not lead to as large a reduction in variance compared to averaging many uncorrelated quantities.
- This means that bagging will not lead to a substantial reduction in variance over a single tree in this scenario.

GT_x

Random Forests

- Random Forests overcome the correlation problem by forcing each split to only consider a subset (typically $m \approx \sqrt{p}$) of the p predictors.
- Therefore, on average $(p - m)/p$ of the splits will not even consider the strong predictor, thereby letting the other predictors have more of a chance.
- This process can be considered to be **decorrelating** the trees, thus making the average of the resulting trees less variable and hence more reliable.
- Note that if $m = p$, then bagging is a special case of random forests.

GT_x

Quiz

1. In Random forest, each tree is built on

- A. All features, all data
- B. Subset of features, all data
- C. Subset of features, subset of data
- D. All features, subset of data

Correct Answer: C

2. Random forest model is interpretable

- A. Yes
- B. No (it's an average of predictions from a number of trees)

Correct Answer: B

GTx

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Boosting

Boosting

- Boosting is another approach to improve prediction results of a decision tree.
- It can also be used for statistical learning methods like regression or classification trees.
- We know how bagging works by making use of bootstrap to create multiple bootstrap data sets and then fitting a separate tree to each bootstrap training set. Then we combine the results of the decision trees to create a single prediction model.
- Boosting works similarly except trees are grown sequentially from previously grown trees but focus their attention on the misclassifications from the prior trees.
- Boosting also involves combining a large number of decision trees, $\hat{f}^1, \dots, \hat{f}^B$

GTx

Boosting for Regression Trees

- The boosting approach learns slowly.
- The construction of each tree depends on the trees that have already been grown.
- Given the current model, we fit a decision tree to the residuals from that model.
- We fit a tree using the current residuals r , rather than the outcome Y , as the response.
- This new decision tree is added into the fitted function in order to update the residuals.
- Each of these trees can be rather small, with just a few terminal nodes, determined by the (interaction depth) parameter d in the algorithm.
- By fitting small trees to the residuals, we slowly improve \hat{f} in areas where it does not perform well.
- The shrinkage parameter λ slows the process even further, allowing more and different shaped trees to attack the residuals.
- Note that we are only covering boosting for regression trees. The details of boosting for classification trees are omitted from the ISLR textbook.

GTx

Algorithm 8.2. Boosting for Regression Trees

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training data set.
2. For $b = 1, 2, \dots, B$, repeat:
 - a) Fit a tree \hat{f}^b with d splits (so it will have $d+1$ terminal nodes) to the training data (X, r) .
 - b) Update \hat{f} by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$
 - c) Update the residuals:

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$$

3. Output the boosted model:

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

GT_x

Boosting's three Tuning Parameters

1. The number of trees B .
 - If B is too large, it could result in overfitting in boosting.
 - Cross-validation is used to select B .
2. λ , the shrinkage parameter
 - A small positive number – usual values are 0.01 or 0.001.
 - Very small λ can require using a large value of B to get good prediction performance
3. The number of splits d , which controls the complexity of the boosted ensemble.
 - $d = 1$ work well, in which case each tree is a **stump** with a single split. This boosted ensemble is fitting an additive model, since each term involves a single variable.
 - More generally, d is the interaction depth and controls the interaction order of the boosted model, since d splits involve at most d variables.

GT_x

Quiz

1. in boosting, because the growth of a particular tree takes into account the other trees that have already been grown, smaller trees are typically sufficient as compared to random forest
 - A. True
 - B. False

Correct Answer: A

2. Smaller trees
 - A. Increase overfitting
 - B. Reduce overfitting

Correct Answer: B

GTx

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Bagging, Random Forests, and Boosting Examples

Bagging and Random Forests on Boston Data Set

- We apply bagging and random forests to the Boston data set, using the **randomForest** package in R.
- The exact results may vary depending on your computer's installed versions of R and randomForest package.
- Remember that bagging is a special case of random forests with $m = p$
- We set up a training and test set for the Boston data set.

GTx

Boston Data Set

```
if (!require(MASS)) install.packages("MASS")
library(MASS)
if (!require(randomForest)) install.packages("randomForest")
library(randomForest)
set.seed (1)
train = sample(1:nrow(Boston), nrow(Boston)/2)
boston.test=Boston[-train , "medv"]
bag.boston=randomForest(medv~., data=Boston, subset=train, mtry=13, importance=TRUE)
bag.boston
```

mtry=13 indicates that all 13 predictors should be considered for each split of the tree, which means that bagging should be done

GTx

Bagging Using the randomForest Package

```
> bag.boston
```

Call:

```
randomForest(formula = medv ~ ., data = Boston, mtry = 13,
importance = TRUE, subset = train)
```

Type of random forest: regression

Number of trees: 500

No. of variables tried at each split: 13

Mean of squared residuals: 14.65748

% Var explained: 83.42


GT_x

Bagging Performance on Test Data Set

```
yhat.bag =
predict(bag.boston,newdata=
Boston[-train,])
plot(yhat.bag, boston.test)
abline(0,1)
```

```
mean((yhat.bag-boston.test)^2)
[1] 13.08278
```

The test set MSE associated with the bagged regression tree is 13.1, almost half that obtained using an optimally-pruned single tree (investigate this on your own).



GT_x

randomForest on Boston Dataset

```
set.seed(1)
rf.boston=randomForest(medv~.,data=Boston,subset=train, mtry=4)
yhat.rf = predict(rf.boston ,newdata=Boston[-train ,])
mean((yhat.rf-boston.test)^2)
```

```
## [1] 13.29981
```

The test MSE using randomForest is 13.3
Compare this MSE to bagging whose MSE was 13.1

GTx

importance() function


```
> importance(rf.boston, type=2)
```

	IncNodePurity
crim	1206.0289
zn	131.2964
indus	961.9492
chas	134.4376
nox	760.9545
rm	5391.3982
age	663.6239
dis	1265.7396
rad	206.9177
tax	475.9189
ptratio	880.9906
black	328.7396
lstat	4909.8992

- The importance function shows the total decrease in node impurity that results from splits over that variable, averaged over all trees.
- In the case of regression trees, the node impurity is measured by the training SSE and for classification trees by the deviance.

GTx

varImpPlot(rf.boston)



Across all the trees considered in the random forest, the two most important variables are:

- **rm** – average number of rooms per dwelling.
- **lstat** - the wealth level of the community as measured by the lower status of the population (percent).

GTx


Boosting on Boston Data Set

- We use the **gbm** package and its `gbm()` function to fit boosted regression trees to the Boston data set
- Run `gbm()` with the option **distribution = "gaussian"** since this is a regression problem; if it were classification then the option would be **distribution = "bernoulli"**
- Set the option **n.trees=5000** to indicate that we want 5000 trees
- Set the option **interaction.depth=4** to limit the depth of each tree

GTx

```
set.seed(1)
boost.boston=gbm(medv~.,data=Boston[train,],distribution="gaussian",
                 n.trees=5000, interaction.depth = 4)
summary(boost.boston)
```

```
> summary(boost.boston)
      var    rel.inf
lstat  lstat 47.1662116
rm      rm  30.8041745
dis     dis  8.1179563
crim    crim 3.4240776
ptratio ptratio 2.0979225
age     age  1.8913236
nox     nox  1.7068723
black   black 1.4067541
chas    chas 1.1739250
tax     tax  0.9062077
indus   indus 0.7930559
rad     rad  0.4543214
```



GTx

Partial Dependence Plots

- lstat and rm are the most important variables
- We can also produce partial dependence plots for these two variables
- These plots produce the marginal effect of the selected variables on the response after integrating out the other variables
- As can be expected, median house prices are increasing with rm and decreasing with lstat


GTx

Partial Dependence Plots

partial dependence plots

```
par(mfrow=c(1,2))
plot(boost.boston, i="rm")
plot(boost.boston, i="lstat")
```

- As can be expected, median house prices are increasing with rm and decreasing with lstat



GTx

Test MSE

```
# Now use the boosted model to predict medv on the test set
> boston.test=Boston[-train,"medv"]
> mean((yhat.boost -boston.test)^2)
[1] 14.65
```

- The test MSE obtained is 14.65.

GTx

Shrinkage Parameter λ Set to 0.02 (Default is .001)

```
boost.boston=gbm(medv~.,data=Boston[train,],distribution="gaussian",
, n.trees=5000, interaction.depth = 4, shrinkage=0.02, verbose=F)
yhat.boost=predict(boost.boston,newdata=Boston[-train,],
n.trees=5000)
mean((yhat.boost -boston.test)^2)
```

```
[1] 17.23
```

- With $\lambda = 0.02$, the test MSE is 17.23

GTx

Recap of Lessons in this Module

- Introduction to Tree Based Methods
- Regression Trees
- Regression Tree Examples
- Classification Trees
- Carseats Example for Classification Trees
- Bagging
- Random Forests
- Boosting
- Bagging, Random Forests, and Boosting Examples

* The materials in this module are from Chapter 8 of the ISLR textbook

GTx

Data Analytics in Business

Classification and Regression Trees (CART)

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

End